

Reg no :j20200776

University ID : w1831982

Pima Indians Diabetes Database

Summary of the dataset:

The data set used for the purpose of this study is Pima Indians Diabetes Database of National Institute of Diabetes and Digestive and Kidney Diseases. This diabetes database, donated by Vincent Sigillito, is a collection of medical diagnostic reports of 768 examples from a population living near Phoenix, Arizona, USA. You can find more information about the dataset

<https://archive.ics.uci.edu/ml/datasets/Pima+Indians+Diabetes> or

<https://www.kaggle.com/uciml/pima-indians-diabetes-database>.

The samples consist of examples with 8 attribute values and one of the two possible outcomes, namely whether the patient is tested positive for diabetes (indicated by output one) or not (indicated by zero).

Exploratory Data Analysis

1. Data loading and cleaning

- Reading data from csv and observing the variables and data

```
# Load the dataset
diabetes_data <- read.csv('diabetes.csv') %>% janitor::clean_names()
```

```
head(diabetes_data) # # visualize the header of Pima data
summary(diabetes_data)
str(diabetes_data)
```

Results:

The screenshot displays the RStudio environment with the following components:

- Source Editor:** Contains R code for loading and cleaning the 'diabetes.csv' dataset. The code includes installing 'corrplot', loading the dataset, visualizing the header, summarizing the data, and creating a cleaned version with 'mutate'.
- Console:** Shows the output of the R commands. It includes the header of the 'diabetes_data' dataset, a summary of the data, and the structure of the 'diabetes_data' object.
- Environment:** Lists the objects in the global environment, including 'correlationMatrix', 'dataset', 'db', 'diabetes_data', 'diabetes_data_cleaned', 'diabetes_data_normaliz...', 'exchangeGBP', 'gdp_exchange_full', 'gdp_exchange_full_data', 'hr_data', 'Ionosphere', 'PimaIndiansDiabetes', 'test_set', 'train_set', 'trainIndex', and 'highlyCorrelated'.

The console output for `summary(diabetes_data)` is as follows:

```
pregnancies glucose blood_pressure skin_thickness insulin bmi diabetes_pedigree_function age outcome
Min. : 0.000 Min. : 0.0 Min. : 0.00 Min. : 0.00 Min. : 0.0 Min. : 0.00 Min. : 0.0780 Min. :21.00 Min. :0.000
1st Qu.: 1.000 1st Qu.: 99.0 1st Qu.: 62.00 1st Qu.: 0.00 1st Qu.: 0.0 1st Qu.:27.30 1st Qu.:0.2437 1st Qu.:24.00 1st Qu.:0.000
Median : 3.000 Median :117.0 Median : 72.00 Median :23.00 Median : 30.5 Median :32.00 Median :0.3725 Median :29.00 Median :0.000
Mean : 3.845 Mean :120.9 Mean : 69.11 Mean :20.54 Mean : 79.8 Mean :31.99 Mean :0.4719 Mean :33.24 Mean :0.349
3rd Qu.: 6.000 3rd Qu.:140.2 3rd Qu.: 88.00 3rd Qu.:32.00 3rd Qu.:127.2 3rd Qu.:36.60 3rd Qu.:0.6262 3rd Qu.:41.00 3rd Qu.:1.000
Max. :17.000 Max. :199.0 Max. :122.00 Max. :99.00 Max. :846.0 Max. :67.10 Max. :2.4200 Max. :81.00 Max. :1.000
```

The console output for `str(diabetes_data)` is as follows:

```
'data.frame': 768 obs. of 9 variables:
 $ pregnancies : int 6 1 8 1 0 5 3 10 2 8 ...
 $ glucose : int 148 85 183 89 137 116 78 115 197 125 ...
 $ blood_pressure : int 72 66 64 66 40 74 50 0 70 96 ...
 $ skin_thickness : int 35 29 0 23 35 0 32 0 45 0 ...
 $ insulin : int 0 0 0 94 168 0 88 0 543 0 ...
 $ bmi : num 33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
 $ diabetes_pedigree_function : num 0.627 0.351 0.672 0.167 2.288 ...
 $ age : int 50 31 32 21 33 30 26 29 53 54 ...
 $ outcome : int 1 0 1 0 1 0 1 1 ...
```

- Here it is notable that it is needed to be normalized as features are in different ranges

```
> str(diabetes_data)
```

```
'data.frame':      768 obs. of  9 variables:
 $ pregnancies      : int  6 1 8 1 0 5 3 10 2 8 ...
 $ glucose          : int  148 85 183 89 137 116 78 115 197 125
 ...
 $ blood_pressure   : int  72 66 64 66 40 74 50 0 70 96 ...
 $ skin_thickness   : int  35 29 0 23 35 0 32 0 45 0 ...
 $ insulin          : int  0 0 0 94 168 0 88 0 543 0 ...
 $ bmi             : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3
30.5 0 ...
 $ diabetes_pedigree_function: num  0.627 0.351 0.672 0.167 2.288 ...
 $ age             : int  50 31 32 21 33 30 26 29 53 54 ...
 $ outcome          : int  1 0 1 0 1 0 1 0 1 1 ...
```

- Here the class variable outcome is needed to be a factor variable.

```
#Need to change the class variable to be a factor variable
diabetes_data$outcome <- factor(diabetes_data$outcome, labels = c("No", "Yes"))
diabetes_data_cleaned =diabetes_data
str(diabetes_data_cleaned)
summary(diabetes_data_cleaned)
```

```
'data.frame':      768 obs. of  9 variables:
 $ pregnancies      : int  6 1 8 1 0 5 3 10 2 8 ...
 $ glucose          : int  148 85 183 89 137 116 78 115 197 125 ...
 $ blood_pressure   : int  72 66 64 66 40 74 50 0 70 96 ...
 $ skin_thickness   : int  35 29 0 23 35 0 32 0 45 0 ...
 $ insulin          : int  0 0 0 94 168 0 88 0 543 0 ...
 $ bmi             : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5
0 ...
 $ diabetes_pedigree_function: num  0.627 0.351 0.672 0.167 2.288 ...
 $ age             : int  50 31 32 21 33 30 26 29 53 54 ...
 $ outcome          : Factor w/ 2 levels "Yes","No": 2 1 2 1 2 1 2 1 2
2 ...
```

```
> summary(diabetes_data_cleaned)
pregnancies      glucose      blood_pressure      skin_thickness
Min.      : 0.000      Min.      : 0.0      Min.      : 0.00      Min.      : 0.00
```

1st Qu.: 1.000	1st Qu.: 99.0	1st Qu.: 62.00	1st Qu.: 0.00
Median : 3.000	Median :117.0	Median : 72.00	Median :23.00
Mean : 3.845	Mean :120.9	Mean : 69.11	Mean :20.54
3rd Qu.: 6.000	3rd Qu.:140.2	3rd Qu.: 80.00	3rd Qu.:32.00
Max. :17.000	Max. :199.0	Max. :122.00	Max. :99.00

skin_thickness	insulin	bmi	diabetes_pedigree_function
Min. : 0.00	Min. : 0.0	Min. : 0.00	Min. :0.0780
1st Qu.: 0.00	1st Qu.: 0.0	1st Qu.:27.30	1st Qu.:0.2437
Median :23.00	Median : 30.5	Median :32.00	Median :0.3725
Mean :20.54	Mean : 79.8	Mean :31.99	Mean :0.4719
3rd Qu.:32.00	3rd Qu.:127.2	3rd Qu.:36.60	3rd Qu.:0.6262
Max. :99.00	Max. :846.0	Max. :67.10	Max. :2.4200

age	outcome
Min. :21.00	No:500
1st Qu.:24.00	Yes :268
Median :29.00	
Mean :33.24	
3rd Qu.:41.00	
Max. :81.00	

- It says that 500 records are mentioned with no diabetes and 268 are having diabetes .
- Agewise analysing against diabetics .

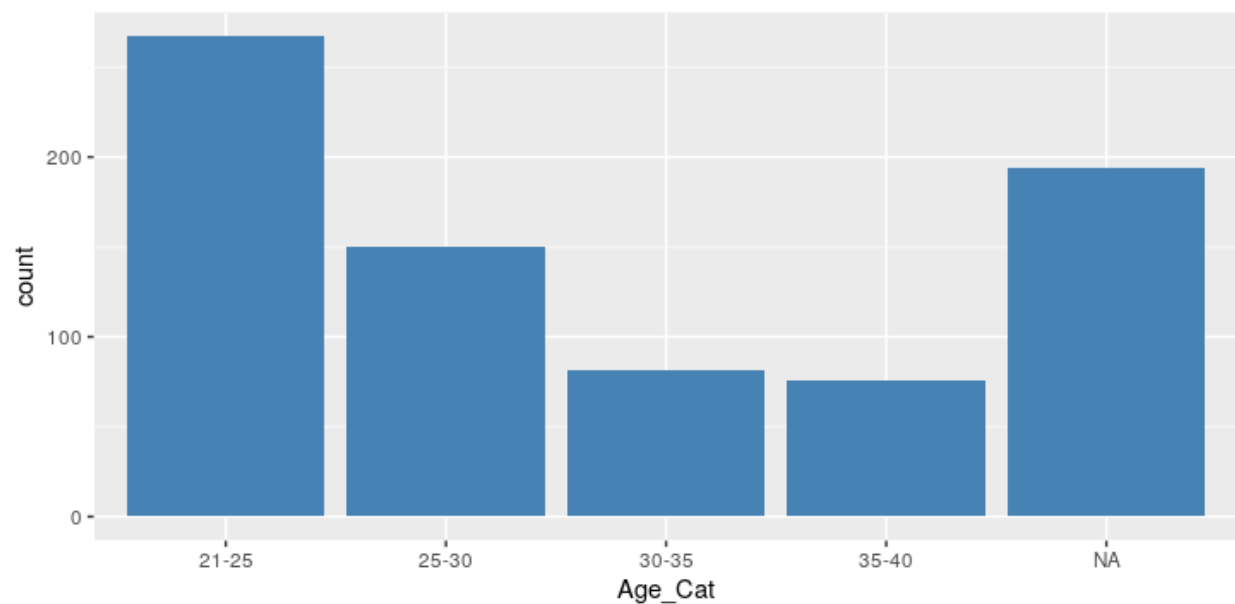

```
##Age waise analysis
db = diabetes_data_cleaned
db$Age_Cat <- ifelse(db$age < 21, "<21",
                    ifelse((db$age>=21) & (db$age<=25), "21-25",
                            ifelse((db$age>25) & (db$age<=30), "25-30",
                                    ifelse((db$age>30) & (db$age<=35),
"30-35",
                                                ifelse((db$age>35) &
(db$age<=40), "35-40",
                                                    ifelse((db$age>40) &
(db$age<=50), "40-50",
```

```

                                                    ifelse((db$age>50)
& (db$age<=60), "50-60", ">60"))))))))

ggplot(aes(x = age), data=db) +
  geom_histogram(binwidth=1, color='black', fill = "#F79420") +
  scale_x_continuous(limits=c(20,90), breaks=seq(20,90,5)) +
  xlab("Age") +
  ylab("No of people by age")
Please refer the image , images/Agwise_analytcis.png


```



Most of the subjects are in between the ages 21 - 30

- check outliers Age_Cat vs BMI

```
library(ggplot2)  
ggplot(aes(x=Age_Cat, y = bmi), data = db) +  
  geom_boxplot() +  
  coord_cartesian(ylim = c(0,70))
```



- Checking the number of missing values in each column.

```
sapply(diabetes_data_cleaned, function(x) sum(is.na(x)))
```

```
pregnancies      glucose      blood_pressure
0               0             0

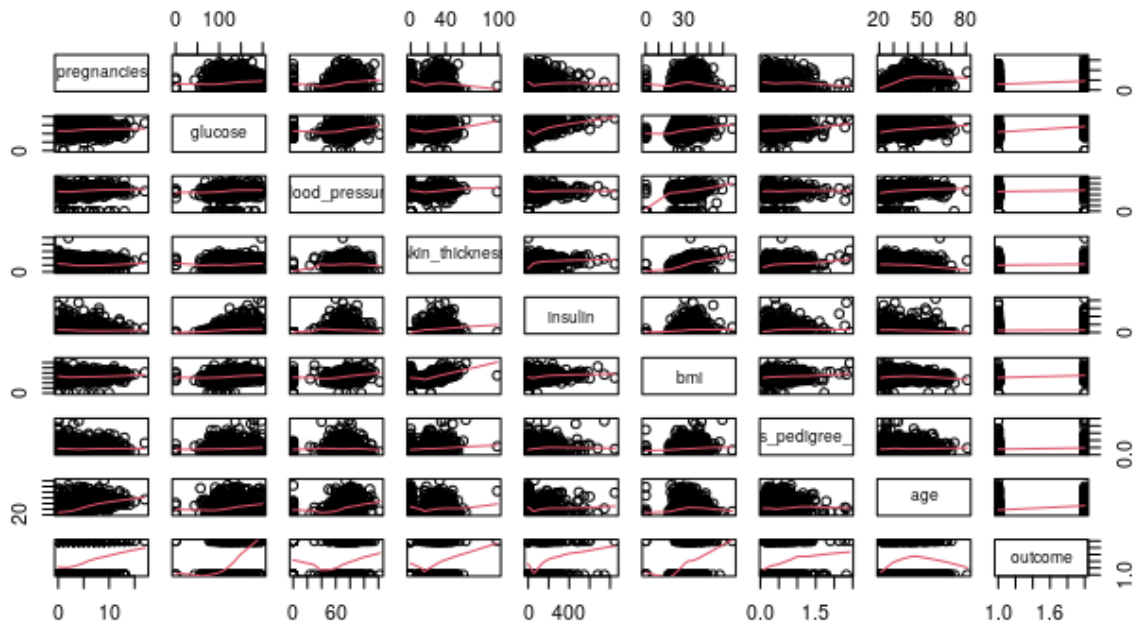
skin_thickness   insulin    bmi      diabetes_pedigree_function
0               0           0             0

age              outcome
0               0
```

- As the results there is no missing values on the data
- Let's produce the matrix of scatterplots

```
pairs(diabetes_data_cleaned, panel = panel.smooth)
```

Please refer the image in the path ,images/matrix_image1.png



- Even explicitly no missing values there are several features that have zero values that are not possible.
- Checking how many zero values are available

```
biological_data_investigation <-
diabetes_data_cleaned[,setdiff(names(diabetes_data_cleaned), c('outcome',
'pregnancies'))]
features_miss_num <- apply(biological_data_investigation, 2, function(x)
sum(x<=0))
features_miss <- names(biological_data_investigation)[ features_miss_num >
0]
Features_miss_num
```

- Result of zero values count are given below

glucose blood_pressure skin_thickness

insulin	bmi	diabetes_pedigree_function
374	11	0
age		
0		

- Investigate how many rows are affected

```
#how many rows are affected
rows_errors <- apply(biological_data_investigation, 1, function(x)
sum(x<=0)>1)
sum(rows_errors)
```

- Results:

234

```
> sum(rows_errors)/nrow(biological_data_investigation)
```

[1] 0.3046875

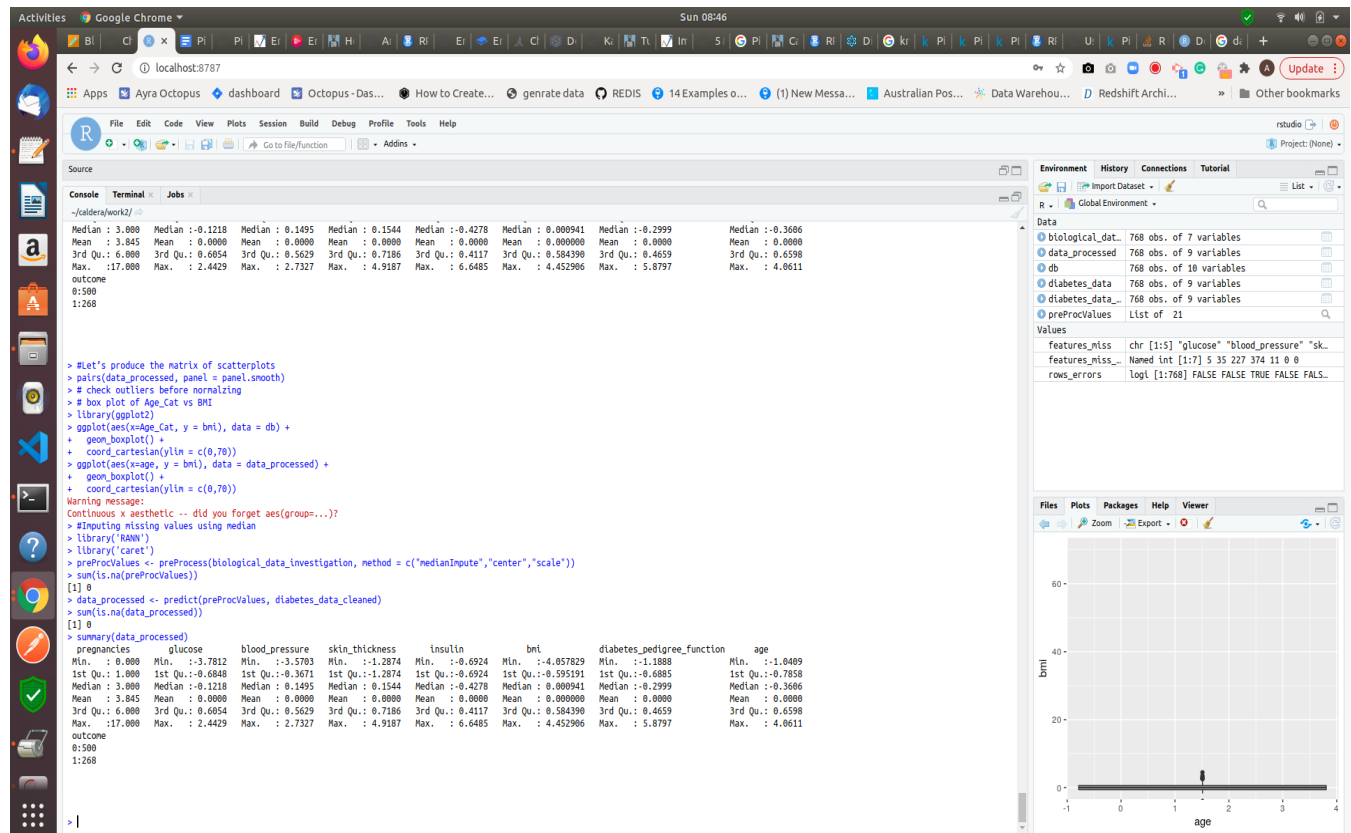
It is more than 30% .

- We are going to try to impute missing data

- we can't remove these rows. We are going to try to impute missing data

```
#Imputing missing values using median
preProcValues <- preProcess(biological_data_investigation, method =
c("medianImpute","center","scale"))
sum(is.na(preProcValues))
```

summary(preProcValues)



pregnancies	glucose	blood_pressure
Min. : 0.000	Min. : -3.7812	Min. : -3.5703
1st Qu.: 1.000	1st Qu.: -0.6848	1st Qu.: -0.3671
Median : 3.000	Median : -0.1218	Median : 0.1495
Mean : 3.845	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 6.000	3rd Qu.: 0.6054	3rd Qu.: 0.5629
Max. : 17.000	Max. : 2.4429	Max. : 2.7327

skin_thickness	insulin	bmi
Min. : -1.2874	Min. : -0.6924	Min. : -4.057829
1st Qu.: -1.2874	1st Qu.: -0.6924	1st Qu.: -0.595191
Median : 0.1544	Median : -0.4278	Median : 0.000941
Mean : 0.0000	Mean : 0.0000	Mean : 0.000000
3rd Qu.: 0.7186	3rd Qu.: 0.4117	3rd Qu.: 0.584390
Max. : 4.9187	Max. : 6.6485	Max. : 4.452906

diabetes_pedigree_function	age	outcome
Min. : -1.1888	Min. : -1.0409	No : 500
1st Qu.: -0.6885	1st Qu.: -0.7858	Yes : 268
Median : -0.2999	Median : -0.3606	
Mean : 0.0000	Mean : 0.0000	
3rd Qu.: 0.4659	3rd Qu.: 0.6598	
Max. : 5.8797	Max. : 4.0611	

2.Variable analysis

- Let's see the proportion of the outcome output .

```
prop.table(table(data_processed$outcome))
```


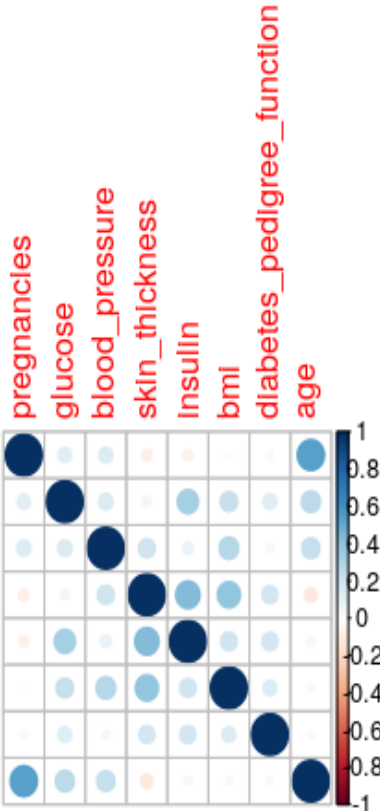
```
No      Yes
0.6510417 0.3489583
```

- Results:

It is quite unbalanced with twice the cases of non diabetes.

- Correlation between variable

```
##correlatinos
# calculate correlation matrix
install.packages("corrplot")
library(corrplot)
correlat <- cor(data_processed[, setdiff(names(data_processed), 'outcome')])
corrplot(correlat)
```



Please refer ,images/matrix_image1.png

```
> print(correlat[,1:3])
```

	pregnancies	glucose	blood_pressure
pregnancies	1.00000000	0.12945867	0.14128198
glucose	0.12945867	1.00000000	0.15258959
blood_pressure	0.14128198	0.15258959	1.00000000
skin_thickness	-0.08167177	0.05732789	0.20737054
insulin	-0.07353461	0.33135711	0.08893338
bmi	0.01768309	0.22107107	0.28180529
diabetes_pedigree_function	-0.03352267	0.13733730	0.04126495
age	0.54434123	0.26351432	0.23952795

- Output of the analysis:

In this case age is highly related with pregnancies (0.55) but it is not more than 75% so no need to eliminate any features.

3.Feature analysis


```
print(importance)
```

ROC curve variable importance

	Importance
glucose	0.7881
bmi	0.6876
age	0.6869
pregnancies	0.6195
diabetes_pedigree_function	0.6062
blood_pressure	0.5865

skin_thickness	0.5536
insulin	0.5379

Please refer the image ,images/feature_analysis.png



It is notable that glucose ,bmi and age contribute high at the sametime insulin contributes low .But here i dont eliminate any features here as my purpose is building ensemble.

4.Boosting algorithm

- Stochastic Gradient Boosting
- C5.0

```
# Create train and test data sets
trainIndex = createDataPartition(data_processed$outcome, p=0.7, list=FALSE)
train_set = data_processed[trainIndex,]
test_set = data_processed[-trainIndex,]
seed <- 10
```

Gradient Boosting

```
fit.gbm <- train(outcome~., data=train_set, method="gbm", metric=metric,
                 trControl=bagcontrol, verbose=FALSE)
```

- Results

```
confusionMatrix(pred_gbm, test_set$outcome)
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	120	18
Yes	30	62

Accuracy : 0.7913
95% CI : (0.733, 0.8419)
No Information Rate : 0.6522
P-Value [Acc > NIR] : 2.878e-06

Kappa : 0.5556

McNemar's Test P-Value : 0.1124

Sensitivity : 0.8000

Specificity : 0.7750

Pos Pred Value : 0.8696

Neg Pred Value : 0.6739

Prevalence : 0.6522

Detection Rate : 0.5217

Detection Prevalence : 0.6000

Balanced Accuracy : 0.7875


'Positive' Class : No

- Analysis of the above results :

Accuracy 79% and Kappa 55%. This is not a bad result.

- ROC

No vs. Yes 0.8625417



After 0.86 sensitivity it gives more accurate results.

- C5.0

```
set.seed(seed)
fit.c50 <- train(outcome~., data=train_set, method="C5.0", metric=metric,
trControl=bagcontrol)

results_boost <- resamples(list( gbm = fit.gbm, c50 = fit.c50))
# Compare models
dotplot(results_boost)
```

Results:

```
> fit.c50
```

C5.0

538 samples
 8 predictor
 2 classes: 'No', 'Yes'

No pre-processing

Resampling: Cross-Validated (10 fold, repeated 3 times)

Summary of sample sizes: 484, 484, 485, 484, 484, 485, ...

Additional sampling using ROSE

Resampling results across tuning parameters:

model	winnow	trials	ROC	Sens	Spec
rules	FALSE	1	0.6923113	0.7647619	0.6058480
rules	FALSE	10	0.7661487	0.7885714	0.6152047
rules	FALSE	20	0.7661487	0.7885714	0.6152047
rules	TRUE	1	0.6920663	0.7600000	0.6022417
rules	TRUE	10	0.7690045	0.7495238	0.6382066
rules	TRUE	20	0.7690045	0.7495238	0.6382066
tree	FALSE	1	0.7359844	0.7380952	0.6475634
tree	FALSE	10	0.7787204	0.7361905	0.6791423
tree	FALSE	20	0.7787204	0.7361905	0.6791423
tree	TRUE	1	0.7236647	0.7457143	0.6352827
tree	TRUE	10	0.7639237	0.7542857	0.6320663
tree	TRUE	20	0.7639237	0.7542857	0.6320663

ROC was used to select the optimal model using the largest value.

The final values used for the model were trials = 10, model = tree and winnow = FALSE.

- Analysis of the above results

It represents that gbm works under this condition.

Comparing both models :

```
results_boost <- resamples(list( gbm = fit.gbm, c50 = fit.c50))  
> # Compare models  
> dotplot(results_boost)  
> results_boost
```

Call:


```
resamples.default(x = list(gbm = fit.gbm, c50 = fit.c50))
```

Models: gbm, c50

Number of resamples: 30

Performance metrics: ROC, Sens, Spec

Time estimates for: everything, final model fit



- Using gbm its possible to boost the accuracy

5. Bagging Algorithms

- Random Forest

Bagging Algorithm (Random Forest)

```
set.seed(seed)
fit.rf <- train(outcome~., data=train_set, method="rf", metric=metric,
trControl=bagcontrol)
# evaluate results on test set
test_set$pred <- predict(fit.rf, newdata=test_set)

#test_set$outcome <- as.factor(test_set$outcome)

str(test_set)
confusionMatrix(data = test_set$pred, test_set$outcome)
pred_fit.rf <- predict(fit.rf, test_set, type="prob")
roc_fit.rf <- roc(test_set$outcome, pred_fit.rf$Yes)
colAUC(pred_fit.rf$Yes, test_set$outcome, plotROC = TRUE)
```

- Result of Random forest bagging

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	118	18
Yes	32	62

Accuracy : 0.7826
95% CI : (0.7236, 0.8341)
No Information Rate : 0.6522
P-Value [Acc > NIR] : 1.156e-05

Kappa : 0.5396

Mcnemar's Test P-Value : 0.06599

Sensitivity : 0.7867
Specificity : 0.7750
Pos Pred Value : 0.8676
Neg Pred Value : 0.6596

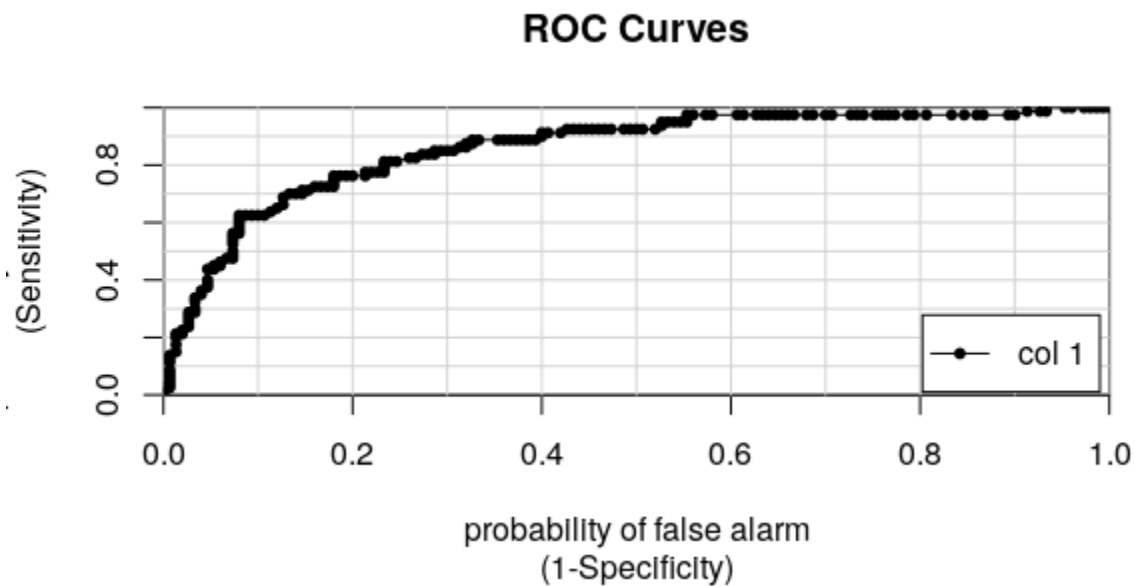
Prevalence : 0.6522
Detection Rate : 0.5130
Detection Prevalence : 0.5913
Balanced Accuracy : 0.7808

'Positive' Class : No

Outcome of the analysis:

78% is accuracy and Kappa is 53% .Most positive to be No value for the class.

No vs. Yes 0.8582917



This says that this model may be biased to predict to be No till improve the model sensitivity (0.85).

- Bagged CART

Bagged CART

```
set.seed(seed)
fit.treebag <- train(outcome~., data=train_set, method="treebag",
metric=metric, trControl=bagcontrol)
pred_fit.treebag <- predict(fit.treebag, newdata=test_set)
confusionMatrix(data = pred_fit.treebag, reference = test_set$outcome)
```

Results:

```
confusionMatrix(data = pred_fit.treebag, reference = test_set$outcome)
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	124	25
Yes	26	55

Accuracy : 0.7783

95% CI : (0.719, 0.8302)

No Information Rate : 0.6522

P-Value [Acc > NIR] : 2.232e-05

Kappa : 0.5127

Mcnemar's Test P-Value : 1

Sensitivity : 0.8267

Specificity : 0.6875

Pos Pred Value : 0.8322

Neg Pred Value : 0.6790

Prevalence : 0.6522

Detection Rate : 0.5391

Detection Prevalence : 0.6478

Balanced Accuracy : 0.7571

'Positive' Class : No

Outcome of the analysis:

77% is accuracy and Kappa is 51% .

When comparing Random forest and Bagged CART then Random forest gives high accuracy.

Finally , Random Forest behaves very well for this scenario .

- Evaluation with test data set using Random forest

evaluate results on test set

```
test_set$pred <- predict(fit.rf, newdata=test_set)
```

```
confusionMatrix(data = test_set$pred, reference = test_set$outcome)
```

Results:

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	118	18
Yes	32	62

Accuracy : 0.7826

95% CI : (0.7236, 0.8341)

No Information Rate : 0.6522

P-Value [Acc > NIR] : 1.156e-05

Kappa : 0.5396

Mcnemar's Test P-Value : 0.06599

Sensitivity : 0.7867

Specificity : 0.7750

Pos Pred Value : 0.8676

Neg Pred Value : 0.6596

Prevalence : 0.6522

Detection Rate : 0.5130

Detection Prevalence : 0.5913

Balanced Accuracy : 0.7808

'Positive' Class : No

- It gives good 78% accuracy results with testing data.

6. Stacking Algorithms

- Classification and Regression Trees (CART),
- K-Nearest Neighbors (KNN),
- Naïve Bayes (NB)

```
# Stacking Algorithms
stack_control <- trainControl(sampling="rose",method="repeatedcv", number=10,
                             repeats=3,
                             savePredictions='final', classProbs=TRUE,
                             summaryFunction = twoClassSummary)
algorithmList <- c( 'knn','rpart','nb')
set.seed(seed)
str(train_set);
#levels(train_set$outcome) <- make.names(levels(factor(train_set$outcome)))
stack_models <- caretList(outcome~., data=train_set, trControl=stack_control,
                          methodList=algorithmList, metric = "ROC" )
stacking_results <- resamples(stack_models)
summary(stacking_results)
dotplot(stacking_results)
names(stack_models)
lapply(stack_models,"[", "results")
# Check correlation between models to ensure the results are uncorrelated and
# can be
modelCor(stacking_results)
splom(stacking_results)
```

- Results:

```
> summary(stacking_results)
```

```
Call:
```

```
summary.resamples(object = stacking_results)
```

```
Models: knn, rpart, nb
```

```
Number of resamples: 30
```

ROC


	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
knn	0.5488722	0.7001984	0.7620301	0.7579588	0.8287594	0.9365079	0
rpart	0.5406015	0.6549499	0.6962406	0.7018101	0.7701128	0.8157895	0
nb	0.6571429	0.7593985	0.7958229	0.7997271	0.8449457	0.9353383	0

Sens

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
knn	0.6000000	0.7714286	0.8285714	0.8095238	0.8571429	0.9142857	0
rpart	0.4571429	0.6357143	0.7142857	0.7133333	0.7714286	1.0000000	0
nb	0.6000000	0.7500000	0.8285714	0.8047619	0.8571429	0.9428571	0

Spec

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
knn	0.3333333	0.4269006	0.5263158	0.5549708	0.6710526	0.9444444	0
rpart	0.3333333	0.5869883	0.6315789	0.6550682	0.7763158	0.8947368	0
nb	0.3333333	0.5336257	0.6052632	0.6005848	0.6842105	0.8421053	0



Confidence level 95%.

Please refer [images/stacking_confidence.png](#)

- Analysis outcomes :
It says the Naive Bayes algorithm works with high accuracy.

- Check correlations between models

```
modelCor(stacking_results)
```

```


knn      rpart      nb
knn      1.0000000 0.5827502 0.7126511
rpart    0.5827502 1.0000000 0.5650957
nb       0.7126511 0.5650957 1.0000000

```

Analysis outcome :

No one is more than 80% so no models are core-related.

- `splom(stacking_results)`



- Stack using naive Bayes

```
# stack using nb
set.seed(seed)
#start time
start_time <- Sys.time()
stack_nb_model <- caretStack(stack_models, method="nb", metric = metric,
                             trControl=stack_control)
end_time <- Sys.time()

timedifferences <- (end_time - start_time)
print(timedifferences)
```

Results:

Time difference of 6.942099 secs

```
confusionMatrix(data = stack.nb.pred, reference = test_set$outcome)
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	128	26
Yes	22	54

Accuracy : 0.7913
95% CI : (0.733, 0.8419)
No Information Rate : 0.6522
P-Value [Acc > NIR] : 2.878e-06


Kappa : 0.5346

Mcnemar's Test P-Value : 0.665

Sensitivity : 0.8533
Specificity : 0.6750
Pos Pred Value : 0.8312
Neg Pred Value : 0.7105
Prevalence : 0.6522
Detection Rate : 0.5565
Detection Prevalence : 0.6696
Balanced Accuracy : 0.7642

'Positive' Class : No

Here Accuracy has been increased to 79% and kappa is 0.5346. This is a good results



No vs. Yes 0.8426667

After 0.6 (in X axis) it gives a more accurate prediction .

- Stack using KNN

```
# stack using knn
set.seed(seed)
#start time
start_time <- Sys.time()
stack_model.knn <- caretStack(stack_models, method="knn", metric = metric,
                             trControl=stack_control)
end_time <- Sys.time()

timedifferences <- (end_time - start_time)
print(timedifferences)
```

Time difference of 2.964403 secs

```
> confusionMatrix(data = pred_fit_stack.knn, reference = test_set$outcome)
```

Confusion Matrix and Statistics

	Reference	
Prediction	No	Yes
No	117	31
Yes	33	49

Accuracy : 0.7217

95% CI : (0.659, 0.7786)

No Information Rate : 0.6522

P-Value [Acc > NIR] : 0.0148

Kappa : 0.3902

Mcnemar's Test P-Value : 0.9005

Sensitivity : 0.7800

Specificity : 0.6125

Pos Pred Value : 0.7905

Neg Pred Value : 0.5976

Prevalence : 0.6522


Detection Rate : 0.5087

Detection Prevalence : 0.6435

Balanced Accuracy : 0.6963

'Positive' Class : No

No vs. Yes 0.7544583



After 0.8 (in X axis) it gives a more accurate prediction .

Here Accuracy has been increased to 72% and kappa is 39%.This is nood when comparing with Naive bayes.

- Stack using CART

```
#start time
start_time <- Sys.time()
stack_rpart_model.rpart <- caretStack(stack_models, method="rpart", metric
= metric,
                                     trControl=stack_control)
end_time <- Sys.time()

timedifferences <- (end_time - start_time)
print(timedifferences)
```

Time difference of 2.542559 secs

```
ConfusionMatrix(data =pred_fit_stack.rpart, reference = test_set$outcome)
```

Confusion Matrix and Statistics

Prediction	Reference	
	No	Yes
No	124	25
Yes	26	55

Accuracy : 0.7783

95% CI : (0.719, 0.8302)

No Information Rate : 0.6522

P-Value [Acc > NIR] : 2.232e-05

Kappa : 0.5127

Mcnemar's Test P-Value : 1

Sensitivity : 0.8267

Specificity : 0.6875

Pos Pred Value : 0.8322

Neg Pred Value : 0.6790

Prevalence : 0.6522

Detection Rate : 0.5391

Detection Prevalence : 0.6478


Balanced Accuracy : 0.7571

'Positive' Class : No

No vs. Yes 0.806125

After 0.85 (in X axis) it gives a more accurate prediction .

Here Accuracy has been increased to 78% and kappa is 51%.This is nood when comparing with Naive bayes.



Compare all three models

When comparing the three stack models Naive Bayes works well 79% and kappa is 0.5346.
Training time comparing :Naive Nayes takes more time ,6.942099 secs

Just single Naive bayes results:

```
#naive bayes
set.seed(seed)
fit.nb <- train(outcome~., data=train_set, method="nb", metric=metric,
trControl=bagcontrol)
pred_fit.nb <- predict(fit.nb, newdata=test_set)
confusionMatrix(data = pred_fit.nb, reference = test_set$outcome)
```


Confusion Matrix and Statistics

Prediction	Reference	
	No	Yes
No	111	23
Yes	39	57

Accuracy : 0.7304

95% CI : (0.6682, 0.7866)

No Information Rate : 0.6522

P-Value [Acc > NIR] : 0.006878

Kappa : 0.4323

Mcnemar's Test P-Value : 0.056780

Sensitivity : 0.7400

Specificity : 0.7125

Pos Pred Value : 0.8284

Neg Pred Value : 0.5938

Prevalence : 0.6522

Detection Rate : 0.4826

Detection Prevalence : 0.5826

Balanced Accuracy : 0.7263

'Positive' Class : No

Without stacking Naive Bayes gives 73% and kappa is 0.43

So it's more clear ensemble stacking gives more accuracy than non-ensembling method of NB .

How can we increase reliability and consistency?

1.As unwanted features were not removed, so if those are removed possible to get more accurate and decrease the training time.

2.Here imbalanced class data has been used ,if we can get balanced class data possible to see more accurately.

3.Need to use proper boosting .

References:

<https://www.javaer101.com/en/article/15507128.html>

<https://www.rdocumentation.org/packages/RANN.L1/versions/2.5.2>